

Sai Anirudh Siddi

Las Vegas, Nevada, United States | (513) 253-7344 | saianirudhsiddi4@gmail.com | [LinkedIn](#) | github.com/Anirudh64210 | www.saianirudh.blog

EDUCATION

University of Cincinnati

Aug 2024 - Dec 2025

Master of Engineering, Computer Science | GPA: 3.95/4.0

Jawaharlal Nehru Technological University, Hyderabad

Jul 2020 - May 2024

Bachelor of Technology, Computer Science | GPA: 3.9/4.0

SKILLS

- **Programming Languages:** Python, SQL, Java, C, C#, Go, Rust, Scala, R, CSS
- **Frameworks & Libraries:** PyTorch, Transformers, LangChain, Scikit-learn, TensorFlow, FastAPI, Node.js, Streamlit
- **AI/ML Systems:** LLMs, Reinforcement Learning (RLHF, RLVR), Reward Modeling, LLM Evaluation, Red-Teaming, Adversarial Prompting, Rubric Design, Mechanistic Interpretability, RAG Pipelines, Agentic AI, FAISS, Machine Learning, MLOps
- **Backend & Infrastructure:** Kubernetes, REST APIs, Microservices, GPU Inference Pipelines, RunPod, Docker, CI/CD Workflows
- **Data Engineering:** Spark, Hadoop, ETL, Pipelines, Validation, Feature Engineering, Pandas, NumPy, Data Management
- **Cloud & Platforms:** AWS, GCP, Azure, Snowflake, Git, Linux, MongoDB, MySQL
- **Monitoring & Observability:** Arize Phoenix, Sentry, Experiment Tracking, Model Evaluation, Analytics Dashboards, Advanced Analytics, Tableau, Power BI, Looker

WORK EXPERIENCE

Handshake AI

Aug 2026 - Present

AI Fellow

- Author adversarial evaluation sets and reinforcement learning environments for frontier LLMs like ChatGPT 5 - Sol and Claude Sonnet, writing 5+ domain-specific scenarios, questions, and rubric-based reward functions to evaluate machine learning model failures, achieving a >50% stumping rate on multi-step quantitative reasoning tasks.
- Build and curate labeled evaluation datasets, applying data management practices and iteratively refining prompts against model responses and reasoning traces to isolate genuine reasoning gaps from prompt ambiguity and specification defects.
- Score and analyze model outputs against calibrated rubrics that serve as RL reward signals, using advanced analytics to surface systematic weakness patterns in multi-step quantitative reasoning that feed back into the next round of environment design.

Great American Insurance Group, Predictive Analytics

May 2025 - Jan 2026

Data Engineer

- Analyzed 1M+ underwriting and litigation records using SQL and Python to apply data science and advanced analytics, surface data inconsistencies and workflow bottlenecks, and improve decision accuracy by 10%.
- Assessed data dependencies across multiple systems and built optimized data management pipelines and reporting solutions, improving data accessibility and operational efficiency.
- Validated an AI-assisted analytics tool with business stakeholders, defining requirements and testing outputs to improve processing efficiency by 60%.

Defense Research and Development of India, Digital Forensics

Aug 2023 - Nov 2023

Software Engineer

- Processed and analyzed 700+ structured and unstructured forensic artifacts per run using advanced analytics to support investigative analysis and data-driven decisions.
- Designed and implemented data ingestion, management, and validation pipelines, ensuring accuracy, consistency, and reliability for downstream analysis.
- Documented data schemas and processing logic while testing system outputs, improving system reliability and investigative efficiency by 30%.

Design Code LLC, IT

May 2022 - Sep 2022

Data Engineer

- Developed SQL and PL/SQL reporting workflows across 20+ enterprise modules, enabling accurate data retrieval and actionable insights for business reporting.
- Analyzed 50K+ monthly transactional records to isolate performance issues and data inconsistencies, improving system performance by 15%.
- Partnered with cross-functional teams to gather requirements and validate reporting outputs against business objectives.

PROJECTS

Toki

On-Device AI Meeting Assistant

- Built Toki, a fully offline meeting intelligence system chaining ASR, speaker diarization, and a 4-bit quantized Qwen3-4B LLM for note generation, running entirely on-device with zero cloud inference cost and no data egress.
- Engineered a GPU inference pipeline for a 0.6B ONNX ASR model and 4B LLM on a 4GB-VRAM consumer GPU using VRAM-aware layer offload and streaming decode, cutting live transcription latency from 21s to ~1s and diarization to 0.014 RTF.
- Owned the full product lifecycle as sole engineer, from model selection and evaluation harnesses to a packaged Windows release validated on real user meetings; achieved 9.2% WER and improved retrieval precision@1 from 82.9% to 87.5%.

GlassBox

Mechanistic Interpretability

- Built GlassBox, a mechinterp system for medical LLMs using Google Gemma 3 4B/2B param, surfacing 15+ SAE feature activations and real-time uncertainty/harmfulness tracking during inference.
- Engineered and operationalized a FastAPI and RunPod GPU microservice pipeline with Claude agents, Arize Phoenix, and Sentry, enabling live interpretability monitoring, model evaluation, and production observability.

ExoSeeker

AI Prediction Platform

- Built ExoSeeker, an AI-powered exoplanet detection platform using Kepler mission data, achieving >90% classification accuracy.
- Designed end-to-end ML pipelines using PyTorch (MLP) and scikit-learn (HistGradientBoosting), including feature engineering, class weighting, hyperparameter tuning, and model validation.
- Won Best Use of NASA Data Award at the NASA Space Apps Challenge (2000+ participants); selected to represent Cincinnati at the Global Challenge.

NeedleHelp

MedTech

- Built an end-to-end IoT + ML software system using Python, C++, and microcontrollers for real-time robotic control.
- Implemented real-time data ingestion, control logic, and feedback loops, achieving >90% system accuracy, integrating hardware and software with fault handling and performance monitoring.
- Won 1st Place (Best Project Overall) at Ohio's largest hackathon, outperforming 140+ teams and 800+ participants.